

CS6200 Information Retrieval

Final Project

- Course Name: Information Retrieval
- Semester: Spring 2015
- Team Members: Wei Xu, Yu Xie, Yue Liu
- Instructor: Prof. Naji

Introduction

Description

This project uses a combination of Java and Python. All retrieval models themselves are written in Java, and everything else including query expansion, evaluation, and other text processing is written in Python.

Contributions

- Wei Xu:
 - Builds BM25 search engine
 - Builds Tf-Idf search engine
 - Documentation
- Yu Xie:
 - Implements query expansions
 - Implements evaluation scripts
 - Documentation
- Yue Liu:
 - Builds Lucene search engine

- Implements stopping parser
- Documentation

Literature and Resources

Retrieval Models

We used 3 retrieval models:

1. BM-25
2. Tf-Idf
3. Lucene (default setting)

BM 25 is the one selected to apply advanced searching options, including stopping, stemming, and query expansion.

Stopping

Stopping is a technique that exclude words that are too common to identify targeted documents when evaluating scores of documents for given queries. The provided [Common Words List](#) is used. When applying stopping, these words are both excluded from documents and queries.

Query Expansion

[Words API](#) is used to retrieve information about synonyms and derivants of English words. It has 2500 free hits each day, which is enough for our project.

Even for search engines that do not apply stopping, we eliminate common words when expanding queries. These words would remain in

queries, but we do not expand them. For example, it doesn't make any sense to expand query "what is cacm" into something like "what which whats is are am be cacm".

We use the same term weights for original query terms and expanded terms, and two expansion techniques are used as follows:

1. Synonym

Expand queries with words that have similar meanings. For example. the word "code" is expanded as 6 other words: "encipher", "cypher", "encrypt" "codification", "inscribe", and "cipher".

2. Derivants

Expand queries with derivants of the original term. For example, the word "code" is expanded as 3 other words: "codify", "coding", and "coder".

External Libraries Required

1. [unirest](#): The HTTP request library;
2. [Natural Language Toolkit](#): The NLP library;

Implementation and Discussion

Other than the relevance information for BM25 model, we did not use any corpus-specific or query-sensitive optimization.

Query Expansion

First, we tokenize the original 64 queries using [parseQueries.py](#).

Then we use [Words API](#) to retrieve ([getWordInfo.py](#)) synonyms and derivants of each term that appear in tokenized queries. After that, we got two files [derivative](#) and [similar](#). The first word of each row of these files comes from original queries, and remaining ones are expanded query terms. When expanding, we exclude common words ([CommonWords.py](#)).

Finally, use [getQueries.py](#) to apply the query expansion. At the same time, we also generate queries with stop words excluded. A combination of stopping and expansion (using derivants) is also applied. Together we got five files in total:

1. [Original Queries](#)
2. [Expanded Queries Using Derivants](#)
3. [Expanded Queries Using Synonyms](#)
4. [Queries without Common Words \(Stopping\)](#)
5. [Expanded Queries Using Derivants without Common Words \(Stopping\)](#)

We use the five queries files above as following: In task 1, we use 1; In task 2, we combine 1 with 2 and 3, respectively; In task 3A, we use 4; In Phase 2 (7th run), we use the combination of 4 and 5.

Retrieval Models

1. Lucene: We are using standard analyzer. To run Lucene, you'll have to include all the files in "lucene" folder to a new project and run "lucene.java"
2. Tf-Idf: For TFIDF, we are using the formula $tf * idf$, where tf , idf are

$$tf_{ik} = \frac{f_{ik}}{\sum_{j=1}^t f_{ij}}$$

$$idf_k = \log \frac{N}{n_k}$$

This part of code has already been included in the "java" folder, just run "Executer.java", you'll get the results in "../results/" folder.

3. BM25: For Parameters, we set $k_1 = 1.2$, $k_2 = 100.0$, $b = 0.75$; And the formula we are using is

$$\sum_{i \in Q} \log \frac{(r_i + 0.5)/(R - r_i + 0.5)}{(n_i - r_i + 0.5)/(N - n_i - R + r_i + 0.5)} \cdot \frac{(k_1 + 1)f_i}{K + f_i} \cdot \frac{(k_2 + 1)qf_i}{k_2 + qf_i}$$

In this search engine, we are utilizing the relevance informations in BM25. More specifically speaking, we retrived all the relevance infomation in the program and calculate N, R, n, r to apply them in our formular above. It's worth mentioned that the filenames in cacm.rel has some mistakes: when the file number is less than four digits, the name will omit the '0' at the frnt digits. To fix this, we mannually add the missing '0's to the relevance document.

To run the BM25 search engine, you'll have to unzip all the files under "java" folder, include all the .java files and execute the "Executer.java" file. The out put will be under "../results/" folder to current "java" folder. Also, please include the jsoup.jar file in the exteral library of the project.

Query-By-Query Analysis for Stemming

Evaluation

All evaluation values are obtained through running [evaluation.py](#).

There are 12 documents that missing relevance information (34, 35, 41, 46, 47, 50, 51, 52, 53, 54, 55, 56). Thus any search engine can never find any relevant documents of these queries, so we exclude them in evaluations. It means all the averaged results are base on the remaining 52 valid queries.

Results

	Mean Average Precision	Mean Reciprocal Rank
Lucene	0.414	0.683
Tf-Idf	0.308	0.521
BM25	0.480	0.732
BM25 (Query Expansion Using Synonyms)	0.465	0.708
BM25 (Query Expansion Using Derivants)	0.468	0.728
BM25 (Stopping)	0.449	0.725
BM25 (Stopping and Derivants)	0.443	0.718

As listed in the table, the performance of BM25 is much better than Lucene and tf-idf models. Beyond the difference among retrieval models, a significant reason is that it takes advantage of the relevance information.

We used two approaches of query expansion. Both synonyms and derivants result in decrease of performance. Synonyms has a worse impact on the results than Derivants. By observing expanded queries, we can see that the synonym API includes too many similar terms, some of which may be irrelevant. For example, it expands "selection" into 7 terms including "option", "survival", "pick", and "choice". Besides species evolution related topic, "selection" and "survival" are not even synonyms. The derivants is less "aggressive" when expanse queries since it only include derivants of terms. Also, the decrease is nonsignificant, so it is difficult to tell which method is better.

In our case, stopping also decreases the performance in both precision and reciprocal rank.

Other evaluation reports including precision at rank 5 and 20, precision and recall of every query at every rank, can be found in directory "evaluation".

Conclusions and Outlook

Relevance information can improve performance of search engines a lot. However these data are usually not available, so we also need to focus on other improvements for search engines.

For query expansion, it is likely that we could over expanse terms. Usually some irrelavant terms are included so that precision decreases. When using query expansion, we need to be very careful and selective for the expanded terms.

Other options like stopping can also have a negative impact.

Bibliography

1. Manning, Christopher D., Prabhakar Raghavan, and Hinrich Schütze. "Introduction to information retrieval/Christopher D." (2008).
2. Croft, W. Bruce, Donald Metzler, and Trevor Strohman. Search engines: Information retrieval in practice. Vol. 283. Reading: Addison-Wesley, 2010.
3. Course Notes and Slides of CS 6200 Spring 2016, Northeastern University.